




# Who 'kill'ed my processes?


Trace Signals through Linux Kernel hacking


# Who am I?

- Leesoo Ahn (lsahn\_at\_ooseel\_dot\_net)
- Software Engineer at Wireless AP company
- Contributor at GNOME Desktop (2018 - 2020)
- ARM64 linux kernel core and virtualization

Signal





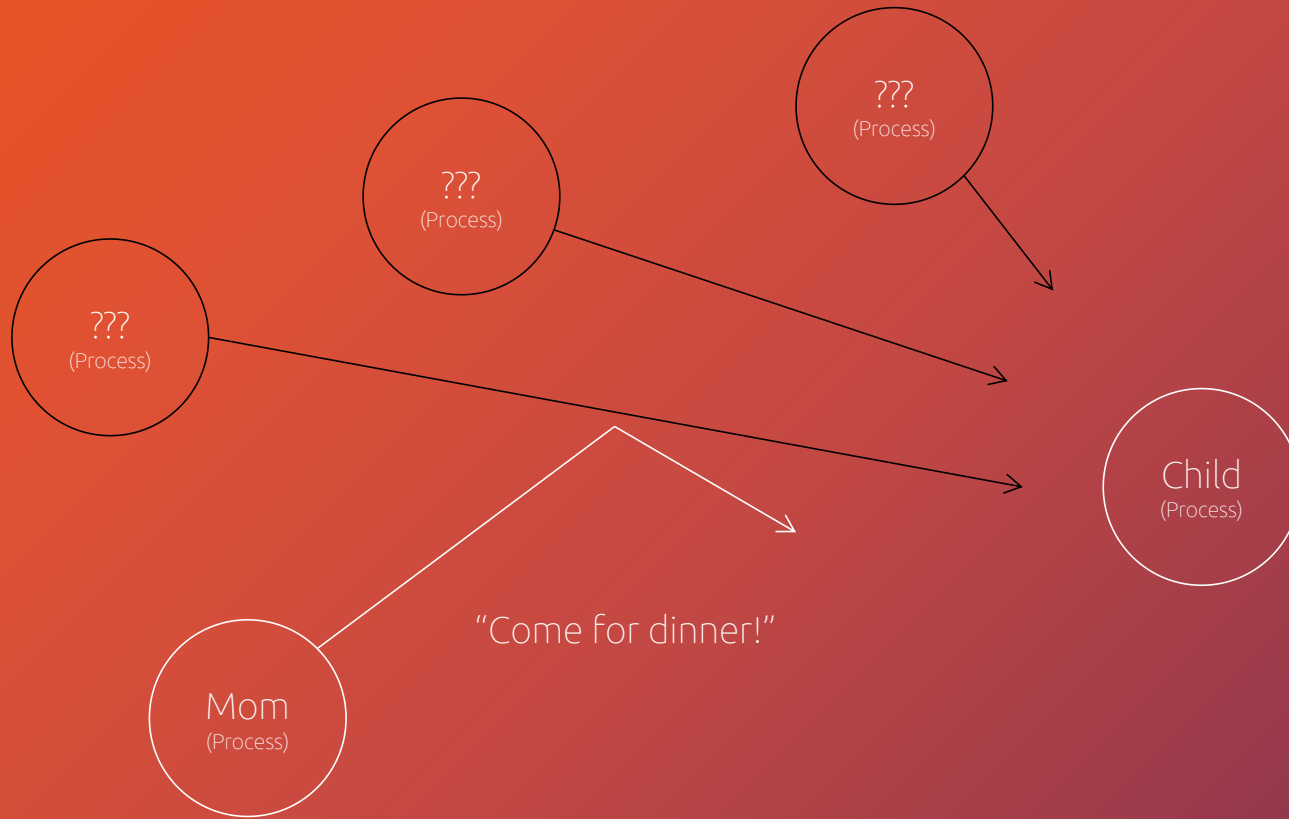




“looks good!”



Real world AIN'T "looks good!"



How real world looks like

We have to survive

We have to survive  
but how?

Some helpers

# The helpers are

- GDB (line by line debugger)

# The helpers are

- GDB (line by line debugger)
- Strace (syscall & signal tracer)

# The helpers are

- GDB (line by line debugger)
- Strace (syscall & signal tracer)
- Ftrace (widely kernel tracer)

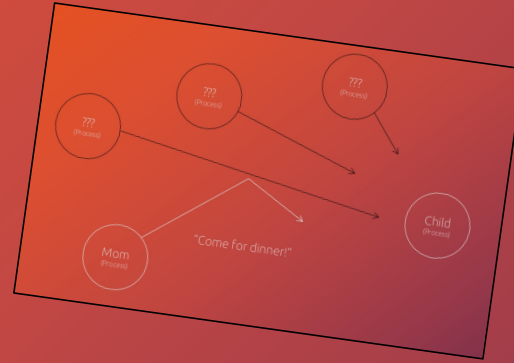


# The helpers are

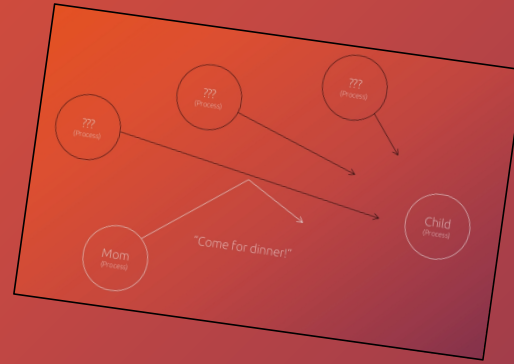
- GDB (line by line debugger)
- Strace (syscall & signal tracer)
- Ftrace (widely kernel tracer)
- et cetera (something I don't know)

What if

What if this



What if this  
happens at booting?



# Oops!

- Mounting a FS which has helpers ATM!

# Oops!

- Mounting a FS which has helpers ATM!
- Deal with order of 'init.d'

# Oops!

- Mounting a FS which has helpers ATM!
- Deal with order of 'init.d'
- Shell ain't allowed

# Oops!


- Mounting a FS which has helpers ATM!
- Deal with order of 'init.d'
- Shell ain't allowed
- et cetera (yet something else)




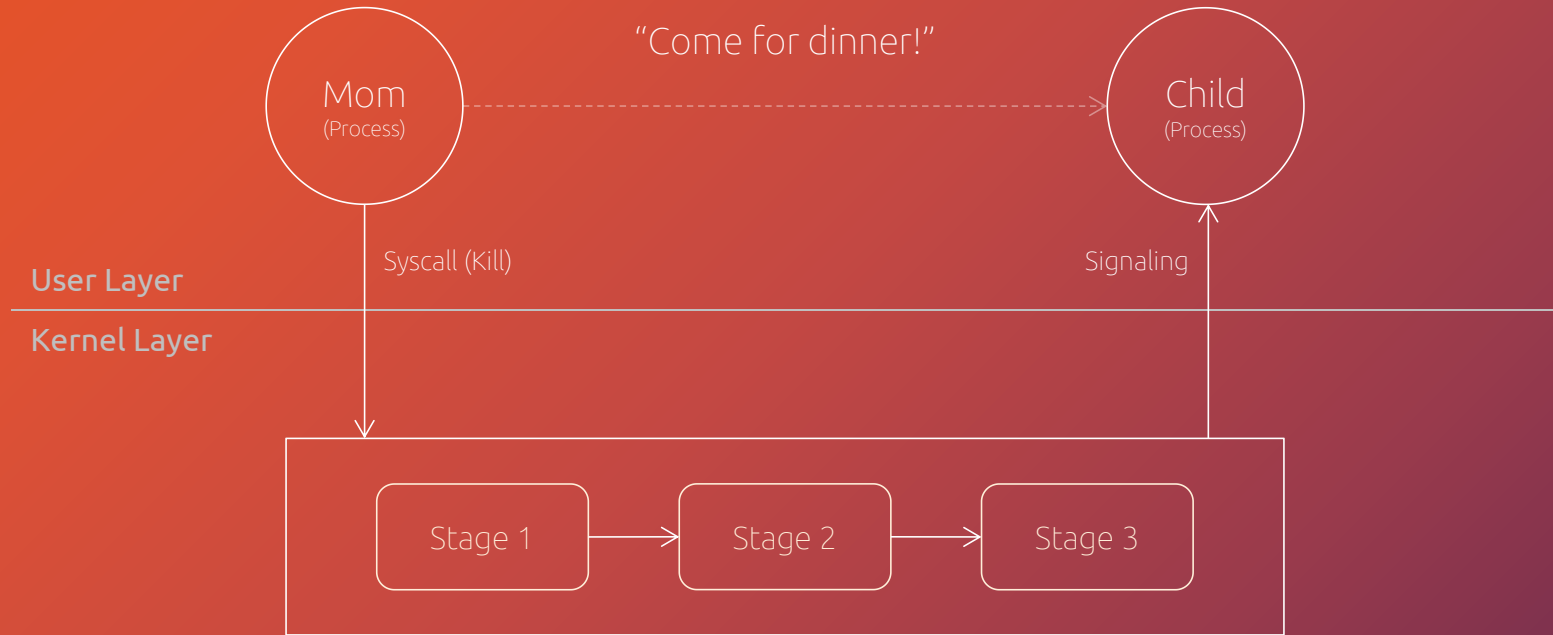
Such a headache!

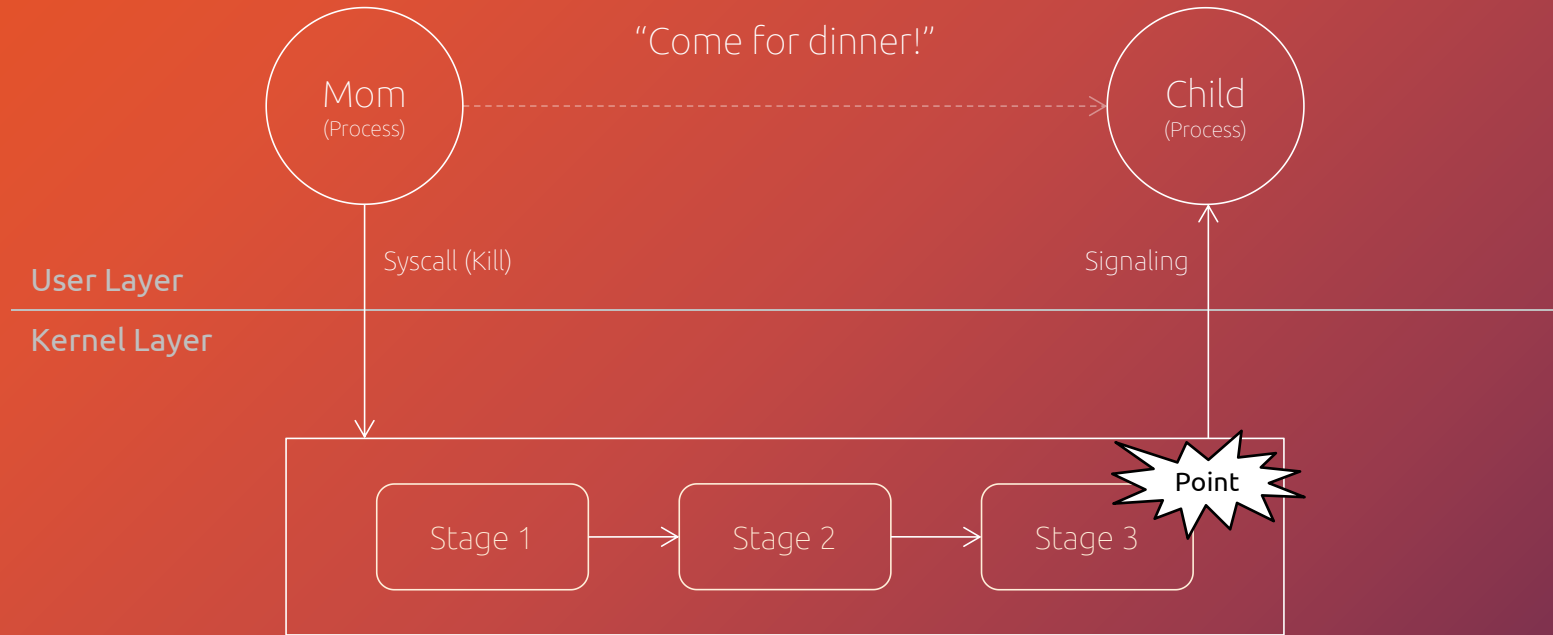
Focus on REAL problem!

What about at 'Kernel'?









## Stage 3, Why?

- Signal Check



# Stage 3, Why?

- Signal Check
- Process Check

## Stage 3, Why?

- Signal Check
- Process Check
- NULL Check

## Stage 3, Why?

- Signal Check
- Process Check
- NULL Check
- and a bunch of Checks!

## Stage 3, How?

- Filter Signals

## Stage 3, How?

- Filter Signals
- Use 'printk'

## Stage 3, How?

- Filter Signals
- Use 'printk'
- Sender, Receiver and Signal number

# Simple Patch

```
complete_signal(sig, t, type);

/* code here */
if (sig != 17 /* SIGCHLD */ &&
    sig != 14 /* SIGALRM */ &&
    info && t)
{
    int srcpid, dstpid;
    char src[TASK_COMM_LEN] = { 0, };
    char dst[TASK_COMM_LEN] = { 0, };
    struct task_struct *cur_task = NULL;

    if (!force) {
        srcpid = info->si_pid;
        cur_task = find_task_by_vpid(srcpid);
    } else
        srcpid = 0;

    dstpid = t->pid;
    memcpy(src, cur_task ? cur_task->comm : "unknown", TASK_COMM_LEN-1);
    memcpy(dst, t->comm, TASK_COMM_LEN-1);

    printk(KERN_INFO "Signal :: (%s %d) --[%d]--> (%s %d)\n",
           src, srcpid, sig, dst, dstpid);
}

ret:
trace_signal_generate(sig, info, t, type != PIDTYPE_PID, result);
return ret;
}
```

- Filter SIGCHLD & SIGALRM
- Sender, Signal and Receiver

# Syscall 'kill' (Kernel 5.4.x)

- SYSCALL\_DEFINE2(kill, ...)
  - kill\_something\_info(...)
  - kill\_pid\_info(...)
  - group\_send\_sig\_info(...)
  - do\_send\_sig\_info(...)
  - send\_signal(...)
  - \_\_send\_signal(...) {  
    complete\_signal(...)  
    /\* code here \*/  
}



# Demo

Thank you!

Q&A